


Impact of Deduplication on F1 Score Performance (Juliet C Dataset with Different Embeddings)



Note: Shaded regions represent 95% confidence intervals (n=20 seeds per k-value). Solid lines: in-distribution test performance. Dashed lines: out-of-distribution (OOD) generalization to Devign dataset.